

central test for logisticR

March 18, 2024

```
[3]: # importing necessary libraries
import pandas as pd
import numpy as np
```

```
[4]: #loading data
data=pd.read_csv("C:\\Users\\SAMUEL K\\Desktop\\banana_quality.csv")
data.head()
```

```
[4]:
```

	Size	Weight	Sweetness	Softness	HarvestTime	Ripeness	Acidity	\
0	-1.924968	0.468078	3.077832	-1.472177	0.294799	2.435570	0.271290	
1	-2.409751	0.486870	0.346921	-2.495099	-0.892213	2.067549	0.307325	
2	-0.357607	1.483176	1.568452	-2.645145	-0.647267	3.090643	1.427322	
3	-0.868524	1.566201	1.889605	-1.273761	-1.006278	1.873001	0.477862	
4	0.651825	1.319199	-0.022459	-1.209709	-1.430692	1.078345	2.812442	

	Quality
0	Good
1	Good
2	Good
3	Good
4	Good

```
[5]: # defining x variables
x=data[['Size', 'Weight', 'Sweetness', 'Softness', 'HarvestTime', 'Ripeness', 'Acidity']]
x.head()
```

```
[5]:
```

	Size	Weight	Sweetness	Softness	HarvestTime	Ripeness	Acidity
0	-1.924968	0.468078	3.077832	-1.472177	0.294799	2.435570	0.271290
1	-2.409751	0.486870	0.346921	-2.495099	-0.892213	2.067549	0.307325
2	-0.357607	1.483176	1.568452	-2.645145	-0.647267	3.090643	1.427322
3	-0.868524	1.566201	1.889605	-1.273761	-1.006278	1.873001	0.477862
4	0.651825	1.319199	-0.022459	-1.209709	-1.430692	1.078345	2.812442

```
[6]: # dropping the unwanted x columns
x=data.drop(columns=["Quality"])
x.head()
```

```
[6]:      Size      Weight  Sweetness  Softness  HarvestTime  Ripeness  Acidity
0 -1.924968  0.468078   3.077832 -1.472177    0.294799  2.435570  0.271290
1 -2.409751  0.486870   0.346921 -2.495099   -0.892213  2.067549  0.307325
2 -0.357607  1.483176   1.568452 -2.645145   -0.647267  3.090643  1.427322
3 -0.868524  1.566201   1.889605 -1.273761   -1.006278  1.873001  0.477862
4  0.651825  1.319199   -0.022459 -1.209709   -1.430692  1.078345  2.812442
```

```
[7]: # defining y variables
y=data["Quality"]
y.head()
```

```
[7]: 0    Good
1    Good
2    Good
3    Good
4    Good
Name: Quality, dtype: object
```

```
[8]: #importing necessary libraries
from sklearn.model_selection import train_test_split,GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import recall_score,accuracy_score,precision_score,f1_score
```

```
[9]: # splitting data
x_train, x_test, y_train, y_test= train_test_split(x,y,test_size=0.2,
↳random_state=42)
```

```
[10]: #building a model
model=LogisticRegression()
model.fit(x_train,y_train)
```

```
[10]: LogisticRegression()
```

```
[11]: #making aprediction for y
pred=model.predict(x_test)
pred
```

```
[11]: array(['Bad', 'Good', 'Good', ..., 'Good', 'Bad', 'Bad'], dtype=object)
```

```
[12]: # calculating for recall score using pos_label
recall=recall_score(y_test,pred ,pos_label="Good")
recall
```

```
[12]: 0.8888888888888888
```

```
[13]: #calculating for accuray score
accuracy=accuracy_score(y_test,pred)
```

```
accuracy
```

```
[13]: 0.879375
```

```
[14]: #calculating precision score  
precision=precision_score(y_test,pred,pos_label="Good")  
precision
```

```
[14]: 0.8771084337349397
```

```
[15]: f1=f1_score(y_test,pred,pos_label="Good")  
f1
```

```
[15]: 0.8829593693147362
```

```
[16]: #standardizing the data  
scaler=StandardScaler()  
x_train_scaled=scaler.fit_transform(x_train)  
x_test_scaled=scaler.transform(x_test)  
#x_train1
```

```
[17]: # define the hyperparameters to tune  
model = LogisticRegression()  
param_grid={'C':[1],  
            'penalty':['l1', 'l2', 'elasticnet', None]  
            }  
param_grid
```

```
[17]: {'C': [1], 'penalty': ['l1', 'l2', 'elasticnet', None]}
```

```
[18]: #performing GridSearch cross validation for hyperparameter  
grid_search= GridSearchCV(model, param_grid,cv=5,n_jobs=-1)  
grid_search
```

```
[18]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,  
               param_grid={'C': [1], 'penalty': ['l1', 'l2', 'elasticnet', None]})
```

```
[19]: #fitting grid search into the model  
import warnings  
warnings.filterwarnings("ignore")  
grid_search.fit(x_train,y_train)
```

```
[19]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,  
               param_grid={'C': [1], 'penalty': ['l1', 'l2', 'elasticnet', None]})
```

```
[20]: # best parameters  
best_params = grid_search.best_params_  
best_params
```

```
[20]: {'C': 1, 'penalty': 'l2'}
```

```
[21]: # fitting the best param with x and y train  
best_model= LogisticRegression(**best_params)  
best_model.fit(x_train,y_train)
```

```
[21]: LogisticRegression(C=1)
```

```
[22]: # predicting for y to view if model can predict the values  
y_pred= best_model.predict(x_test)  
y_pred
```

```
[22]: array(['Bad', 'Good', 'Good', ..., 'Good', 'Bad', 'Bad'], dtype=object)
```

```
[24]: #calculating for accuray score  
accuracy2=accuracy_score(y_test,y_pred)  
accuracy2
```

```
[24]: 0.879375
```

```
[26]: if accuracy2> accuracy:  
      print("The optimized model better than ordinary model")  
  
      elif accuracy2<accuracy:  
          print("The ordinary model is greater than optimized model")  
      else:  
          print("Both models have the same accuracy")
```

Both models have the same accuracy

```
[ ]:
```